

1. (a) when  $l(f(x), y) = (f(x) - y)^2$

$$E[l(f(x), y)] = \int_x \int_y (f(x) - y)^2 P(x, y) dy dx$$

$$= \int_x \int_y (f(x) - E(y|x) + E(y|x) - y)^2 \cdot P(x, y) dy dx$$

$$= \int_x \int_y [(f(x) - E(y|x))^2 + (y - E(y|x))^2 - 2(f(x) - E(y|x))(y - E(y|x))] P(x, y) dy dx$$

Let's look at the last term:

$$\int_x \int_y (f(x) - E(y|x))(y - E(y|x)) P(x, y) dx dy$$

$$= E(f(x)) \cdot E(y) - E(f(x) \cdot E(y)) - E(y) \cdot E(y) + E(y^2)$$

$$= 0$$

Since we can not modify  $y$  or  $E(y|x)$ , In order to minimize  $E[l(f(x), y)]$   
 $f(x) = E(y|x)$ , Conditional Expectation of  $y$

(b) when  $l(f(x), y) = |f(x) - y|$

$$E[l(f(x), y)] = \int_x \int_y |f(x) - y| P(x, y) dx dy$$

$$= \int_x \underbrace{\left( \int_y |f(x) - y| P(y|x) dy \right)}_{\text{bracket}} \cdot P(x) dx$$

Because  $x$  can be chosen independent of  $f(x)$ , we can ignore integral over  $x$  and instead focus on how to minimize curly brace

If we take derivative and set to 0:

$$\frac{\partial}{\partial f(x)} \int_y |f(x) - y| P(y|x) dy$$

$$= \frac{\partial}{\partial f(x)} \left[ \int_{-\infty}^{f(x)} (f(x) - y) P(y|x) dy + \int_{f(x)}^{+\infty} (y - f(x)) P(y|x) dy \right] = 0$$

$$\Rightarrow \int_{-\infty}^{f(x)} P(y|x) dy = \int_{f(x)}^{+\infty} P(y|x) dy$$

Optimal  $f(x)$  is the conditional median



2. Correct

Let us consider a classifier  $f \neq f^*$

$$P(f \neq y|x)$$

$$= 1 - P(f=1 \& y=1|x) - P(f=-1 \& y=-1|x)$$

$$= 1 - P(f=1|y=1, x) \cdot P(y=1|x) - P(f=-1|y=-1, x) \cdot P(y=-1|x)$$

Because for each  $(x, y)$ ,  $f(x)$ 's value can be chosen independently of  $y$

$$= 1 - P(f=1|x) \cdot P(y=1|x) - P(f=-1|x) \cdot P(y=-1|x)$$

Because  $(x, y)$  is drawn from fixed distribution  $D$ , Let  $\alpha(x)$  be  $\alpha(x) = P(y=1|x) = 1 - P(y=-1|x)$

$$P(f \neq y|x) = 1 - \mathbb{1}(f=1) \cdot \alpha(x) - \mathbb{1}(f=-1) \cdot (1 - \alpha(x))$$

$$P(f \neq y|x) - P(f^* \neq y|x)$$

$$= \alpha(x) \cdot [\mathbb{1}(f^*=1) - \mathbb{1}(f=1)] + (1 - \alpha(x)) [\mathbb{1}(f^*=-1) - \mathbb{1}(f=-1)]$$

$$\text{Note } \mathbb{1}(f^*=-1) = 1 - \mathbb{1}(f^*=1), \mathbb{1}(f=-1) = 1 - \mathbb{1}(f=1)$$

$$= \alpha(x) \cdot [\mathbb{1}(f^*=1) - \mathbb{1}(f=1)] - (1 - \alpha(x)) [\mathbb{1}(f^*=1) - \mathbb{1}(f=1)]$$

$$= (2\alpha(x) - 1) [\mathbb{1}(f^*=1) - \mathbb{1}(f=1)]$$

$$\text{Note } f^*(x) = \begin{cases} 1 & \text{if } \alpha(x) \geq \frac{1}{2} \\ -1 & \text{else} \end{cases}$$

For  $x$  such that  $\alpha(x) \geq \frac{1}{2}$ ,

$$P(f \neq y|x) - P(f^* \neq y|x) = (2\alpha(x) - 1) [\underbrace{\mathbb{1}(f^*=1)}_{=1} - \underbrace{\mathbb{1}(f=1)}_{=0 \text{ or } 1}]$$

$\geq 0 \quad \geq 0 \quad \geq 0$

For  $x$  such that  $\alpha(x) < \frac{1}{2}$

$$P(f \neq y|x) - P(f^* \neq y|x) = (2\alpha(x) - 1) [\underbrace{\mathbb{1}(f^*=1)}_{=0} - \underbrace{\mathbb{1}(f=1)}_{=0 \text{ or } 1}]$$

$\geq 0 \quad < 0 \quad \leq 0$

Thus  $P(f \neq y|x) - P(f^* \neq y|x) \geq 0$  always hold

Conclusion  $L(f^*) \leq L(f)$  proved.

3:

(i)

First calculate between class  $S_b$ , within class  $S_w$

$$S_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T$$

$$S_w = \sum_{\mathbf{x}_n \in C_1} (\mathbf{x}_n - \mathbf{m}_1)(\mathbf{x}_n - \mathbf{m}_1)^T + \sum_{\mathbf{x}_n \in C_2} (\mathbf{x}_n - \mathbf{m}_2)(\mathbf{x}_n - \mathbf{m}_2)^T$$

Then

calculate the largest absolute eigenvalue as well as corresponding eigenvectors for

$\text{inv}(S_w) * S_b$

Use this eigenvector to project  $X$  and discriminant value by left multiply  $w^T$ , and discriminant value is set as the average of two class mean vectors. And then compare the projected  $X$  value and project discriminant value.

train error rate: 0.310283625731

test error rate: 0.310785714286

test error standard deviation: 0.220061320304

We can not project Boston50 dataset to 2D subspace, because the  $S_B$  is composed of sum of  $K$  matrices, each of which is an outer product of two vectors and thus  $S_B$  has rank 1. In addition, only  $(K-1)$  matrices are independent due to constraint that overall mean = weighted average of each class. So  $S_B$  has at most  $(K-1)$  rank and at most  $(K-1)$  nonzero eigenvalues. Because we need projection matrix to project onto subspace, we can not find more than  $(K-1)$  eigenvectors to compose the projection matrix and can not find more than  $(K-1)$  dimensional subspace.

(ii)

Do the same calculations of  $S_b$  and  $S_w$ , except the  $S_w$  consists of sum of ten classes.

Because we want to project to 2D dimensional subspace, we choose two eigenvectors with largest eigenvalues.

Stack two eigenvectors horizontally as projection matrix to project whole dataset  $X$

In 2D subspace, calculate the covariance matrix and mean using bi\_variate joint density function

$$p(\mathbf{x}|C_k) = \frac{1}{(2\pi)^{d/2}} \frac{1}{|\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2}(\mathbf{x} - \mu_k)^T \Sigma^{-1}(\mathbf{x} - \mu_k) \right\}$$

$P(C_k|x)$  is proportional to  $P(x|C_k) * P(C_k)$

Prior  $P(C_k) = N_k/N$

After we calculate covariance and mean, we can use them to predict  $P(C_k|x)$  and choose the  $k$  corresponding to the largest posterior probability.

train error rate: 0.817658457499

test error rate: 0.830756292425

test error standard deviation: 0.0421947670347

4.

Summary of algorithm

(i) Logistic regression:

First, I will check the number of class  $K$  is larger than 2 or not.

I used the iterative update approach called gradient descent:

Following is for two class classification ( $K=2$ ):

$y = P(C|x) = 1/(1 + \exp(-w^T x + w_0))$ , we use sigmoid function to calculate  $y$

```
For  $j = 0, \dots, d$ 
   $w_j \leftarrow \text{rand}(-0.01, 0.01)$ 
Repeat
  For  $j = 0, \dots, d$ 
     $\Delta w_j \leftarrow 0$ 
  For  $t = 1, \dots, N$ 
     $o \leftarrow 0$ 
    For  $j = 0, \dots, d$ 
       $o \leftarrow o + w_j x_j^t$ 
     $y \leftarrow \text{sigmoid}(o)$ 
    For  $j = 0, \dots, d$ 
       $\Delta w_j \leftarrow \Delta w_j + (r^t - y) x_j^t$ 
    For  $j = 0, \dots, d$ 
       $w_j \leftarrow w_j + \eta \Delta w_j$ 
Until convergence
```

If the classification problem is a multi-class classification ( $K > 2$ ), use following update rule: We use softmax function to calculate  $y$ .

$$y = \exp(o_k) / (\sum_i \exp(o_i))$$

$$o_k = w_k^T X$$

```

For  $i = 1, \dots, K$ 
  For  $j = 0, \dots, d$ 
     $w_{ij} \leftarrow \text{rand}(-0.01, 0.01)$ 
Repeat
  For  $i = 1, \dots, K$ 
    For  $j = 0, \dots, d$ 
       $\Delta w_{ij} \leftarrow 0$ 
  For  $t = 1, \dots, N$ 
    For  $i = 1, \dots, K$ 
       $o_i \leftarrow 0$ 
      For  $j = 0, \dots, d$ 
         $o_i \leftarrow o_i + w_{ij} x_j^t$ 
      For  $i = 1, \dots, K$ 
         $y_i \leftarrow \exp(o_i) / \sum_k \exp(o_k)$ 
      For  $i = 1, \dots, K$ 
        For  $j = 0, \dots, d$ 
           $\Delta w_{ij} \leftarrow \Delta w_{ij} + (r_i^t - y_i) x_j^t$ 
  For  $i = 1, \dots, K$ 
    For  $j = 0, \dots, d$ 
       $w_{ij} \leftarrow w_{ij} + \eta \Delta w_{ij}$ 
Until convergence

```

Both these two rules came from Introduction to Machine Learning book.

(ii)

Naive bayes with marginal Gaussian distribution:

$P(\mathbf{x}|\mathbf{C}_k) = P(x_1|\mathbf{C}_k) \cdot P(x_2|\mathbf{C}_k) \dots \cdot P(x_d|\mathbf{C}_k) \cdot P(\mathbf{C}_k)$

$$p(\mathbf{x}|\mathbf{C}_k) = \prod_{i=1}^D p(x_i|\mathbf{C}_k) = \frac{1}{(2\pi)^{D/2} \left( \prod_{i=1}^D \sigma_{ik} \right)} \exp \left\{ - \sum_{i=1}^D \frac{(x_i - \mu_{ik})^2}{2\sigma_{ik}^2} \right\}$$

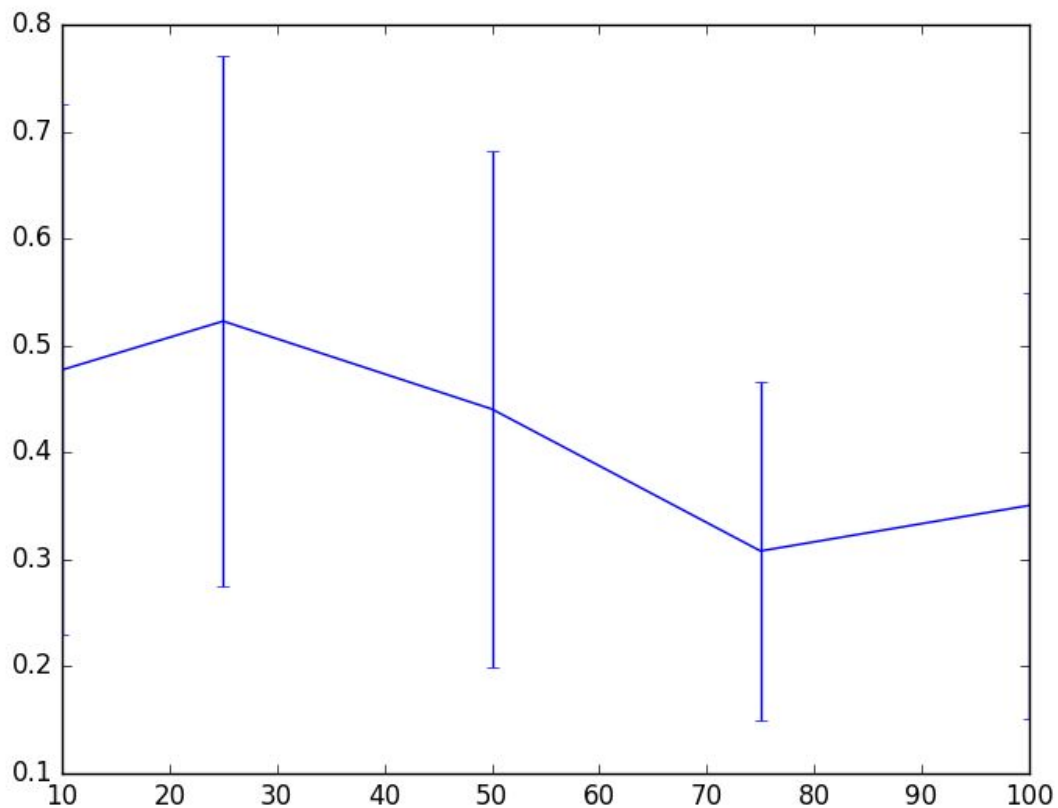
$P(\mathbf{C}_k) = N_k/N$

First, we calculate the sigma, mean and prior probability.

Then we use  $\mathbf{X}$  to calculate posterior probabilities of different classes and select the largest one as class label

Boston50 dataset:

(i)



Mean\_error\_rate:

[ 0.47722772 0.52277228 0.44059406 0.30792079 0.35049505]

Error\_matrix:(row index represents num\_split and column index is train\_percent) :

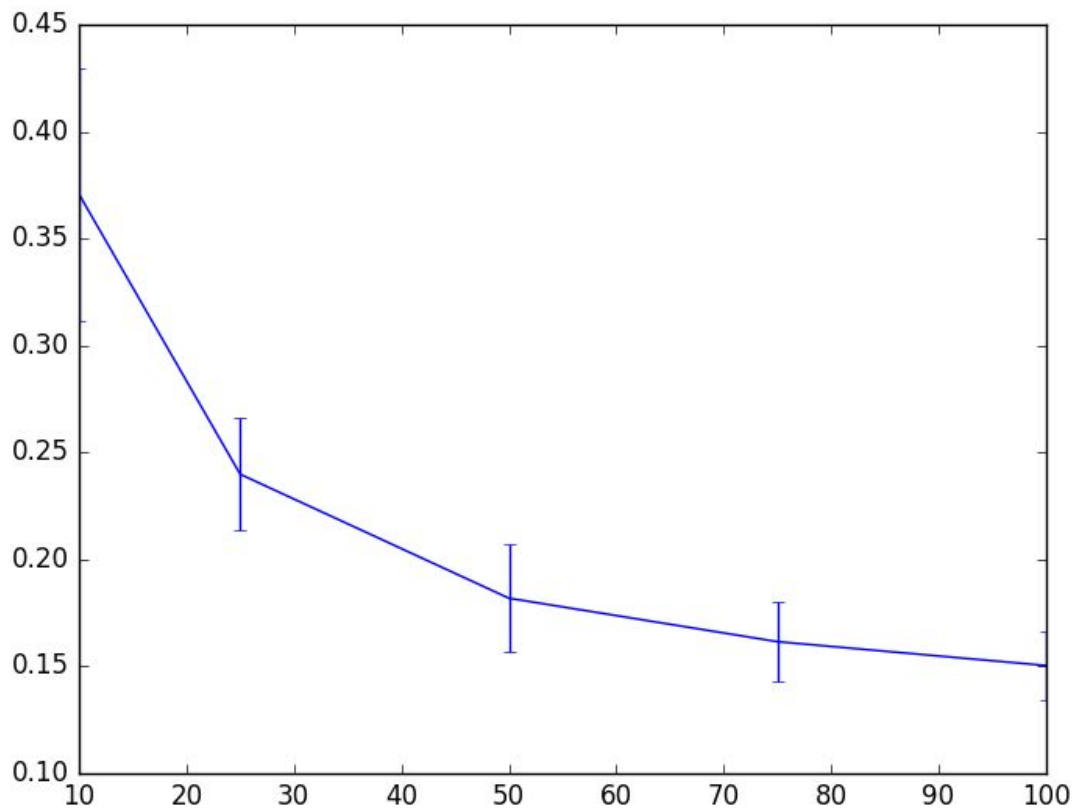
```
[[ 0.54455446 0.45544554 0.45544554 0.45544554 0.54455446]
 [ 0.10891089 0.89108911 0.89108911 0.10891089 0.10891089]
 [ 0.85148515 0.14851485 0.14851485 0.14851485 0.14851485]
 [ 0.22772277 0.77227723 0.77227723 0.22772277 0.22772277]
 [ 0.08910891 0.91089109 0.08910891 0.08910891 0.08910891]
 [ 0.47524752 0.52475248 0.52475248 0.52475248 0.47524752]
 [ 0.57425743 0.42574257 0.42574257 0.42574257 0.57425743]
 [ 0.61386139 0.38613861 0.38613861 0.38613861 0.61386139]
 [ 0.78217822 0.21782178 0.21782178 0.21782178 0.21782178]
 [ 0.5049505 0.4950495 0.4950495 0.4950495 0.5049505 ]]
```

Error standard deviation:

[ 0.24799164 0.24799164 0.24184574 0.15850554 0.19916501]

(ii)





Mean\_error\_rate:

[ 0.34653465 0.23861386 0.20792079 0.1960396 0.18613861]

Error\_matrix:(row index represents num\_split and column index is train\_percent)

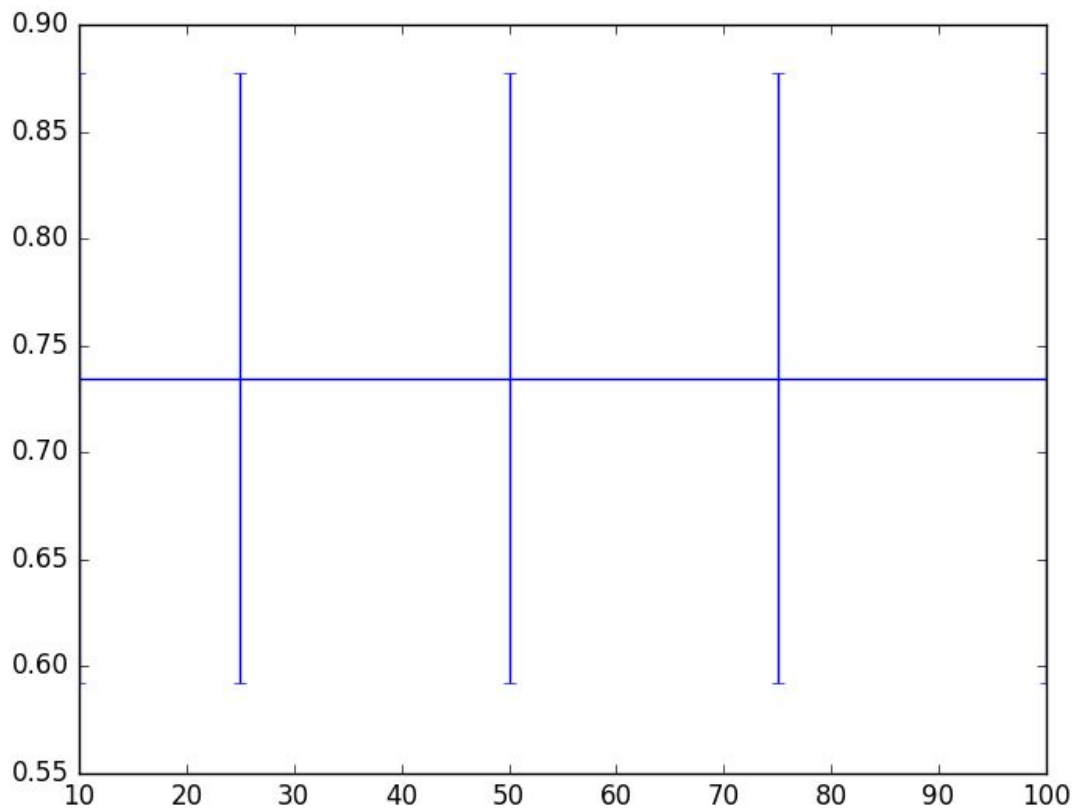
```
[[ 0.4950495 0.14851485 0.27722772 0.16831683 0.10891089]
 [ 0.5049505 0.05940594 0.20792079 0.10891089 0.04950495]
 [ 0.17821782 0.27722772 0.18811881 0.16831683 0.13861386]
 [ 0.38613861 0.34653465 0.26732673 0.3960396 0.3960396 ]
 [ 0.48514851 0.24752475 0.21782178 0.26732673 0.27722772]
 [ 0.14851485 0.22772277 0.0990099 0.13861386 0.14851485]
 [ 0.45544554 0.27722772 0.24752475 0.14851485 0.18811881]
 [ 0.13861386 0.27722772 0.16831683 0.12871287 0.13861386]
 [ 0.4950495 0.24752475 0.21782178 0.26732673 0.27722772]
 [ 0.17821782 0.27722772 0.18811881 0.16831683 0.13861386]]
```

Error standard deviation:

[ 0.15522782 0.07623762 0.04930653 0.08375559 0.09668559]

Boston75 dataset:

(i)



Mean\_error\_rate:

[ 0.73465347 0.73465347 0.73465347 0.73465347 0.73465347]

Error\_matrix:(row index represents num\_split and column index is train\_percent)

[[ 0.87128713 0.87128713 0.87128713 0.87128713 0.87128713]

[ 0.3960396 0.3960396 0.3960396 0.3960396 0.3960396 ]

[ 0.81188119 0.81188119 0.81188119 0.81188119 0.81188119]

[ 0.58415842 0.58415842 0.58415842 0.58415842 0.58415842]

[ 0.82178218 0.82178218 0.82178218 0.82178218 0.82178218]

[ 0.67326733 0.67326733 0.67326733 0.67326733 0.67326733]

[ 0.7029703 0.7029703 0.7029703 0.7029703 0.7029703 ]

[ 0.83168317 0.83168317 0.83168317 0.83168317 0.83168317]

[ 0.79207921 0.79207921 0.79207921 0.79207921 0.79207921]

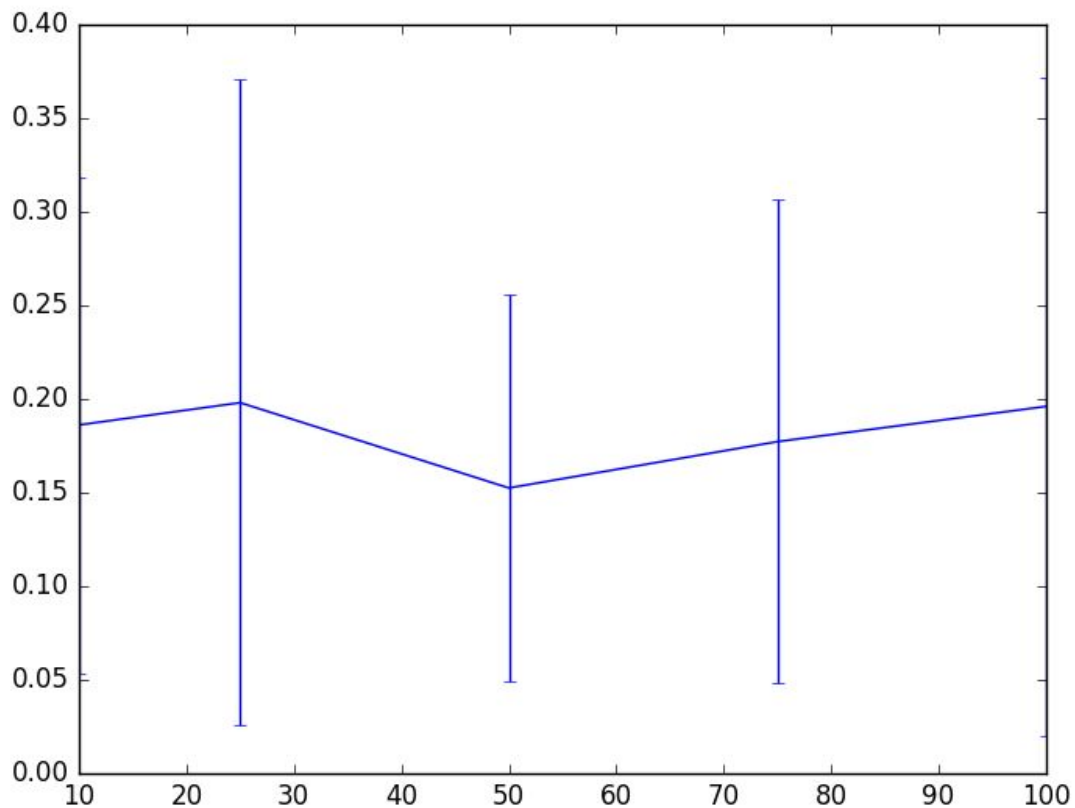
[ 0.86138614 0.86138614 0.86138614 0.86138614 0.86138614]]

Error standard deviation:

[ 0.14250548 0.14250548 0.14250548 0.14250548 0.14250548]

(ii)





Mean\_error\_rate:

[ 0.18613861 0.1980198 0.15247525 0.17722772 0.1960396 ]

Error\_matrix:(row index represents num\_split and column index is train\_percent)

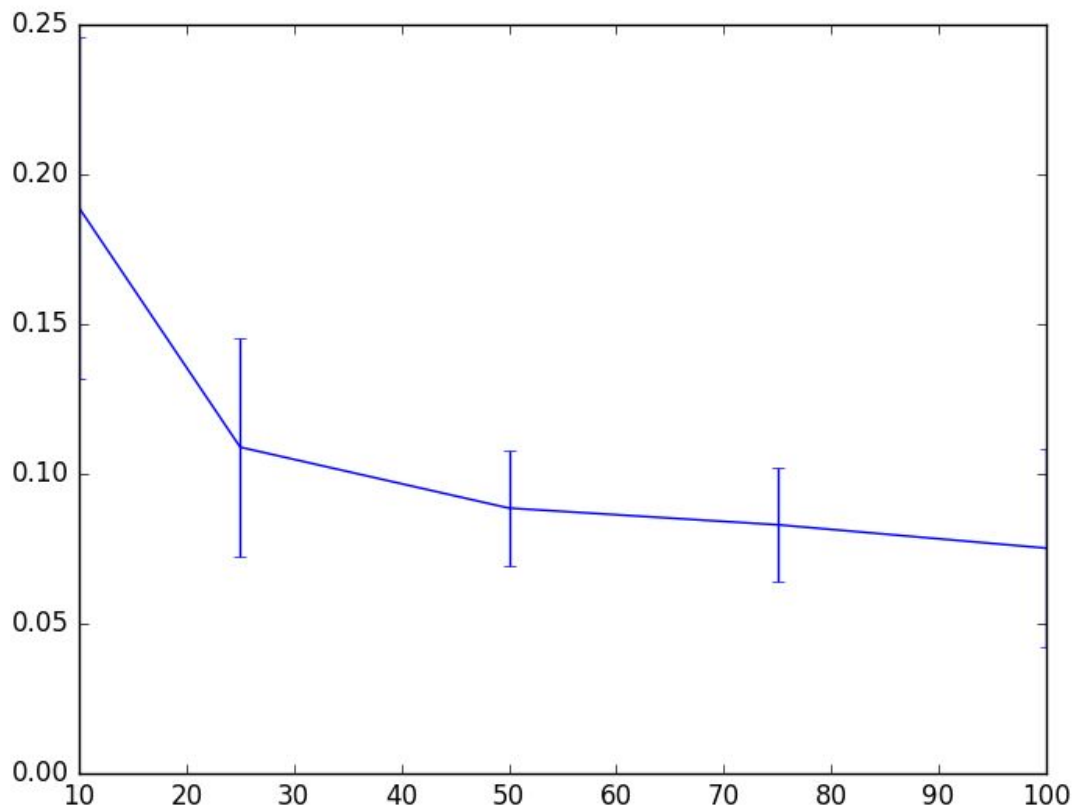
```
[[ 0.23762376 0.10891089 0.11881188 0.0990099 0.08910891]
 [ 0.3960396 0.35643564 0.25742574 0.21782178 0.30693069]
 [ 0.03960396 0.03960396 0.04950495 0.06930693 0.03960396]
 [ 0.02970297 0.02970297 0.02970297 0.04950495 0.02970297]
 [ 0.30693069 0.41584158 0.30693069 0.41584158 0.48514851]
 [ 0.35643564 0.21782178 0.14851485 0.16831683 0.20792079]
 [ 0.14851485 0.15841584 0.17821782 0.16831683 0.16831683]
 [ 0.23762376 0.54455446 0.31683168 0.41584158 0.52475248]
 [ 0.06930693 0.06930693 0.06930693 0.07920792 0.06930693]
 [ 0.03960396 0.03960396 0.04950495 0.08910891 0.03960396]]
```

Error standard deviation:

[ 0.13245142 0.17251605 0.10338834 0.12927903 0.17576977]

Digits dataset:

(i)



Mean\_error\_rate:

[ 0.18885794 0.10891365 0.08857939 0.08300836 0.07520891]

Error\_matrix:(row index represents num\_split and column index is train\_percent)

[[ 0.13649025 0.06128134 0.07520891 0.07799443 0.05292479]

[ 0.28412256 0.15320334 0.11977716 0.09749304 0.11420613]

[ 0.15320334 0.09470752 0.07799443 0.08913649 0.06406685]

[ 0.15320334 0.06963788 0.06685237 0.06685237 0.02785515]

[ 0.13091922 0.07799443 0.06685237 0.05571031 0.04456825]

[ 0.27019499 0.15598886 0.09192201 0.10584958 0.11420613]

[ 0.2005571 0.1281337 0.11699164 0.11142061 0.10027855]

[ 0.14206128 0.08635097 0.0724234 0.05292479 0.0362117 ]

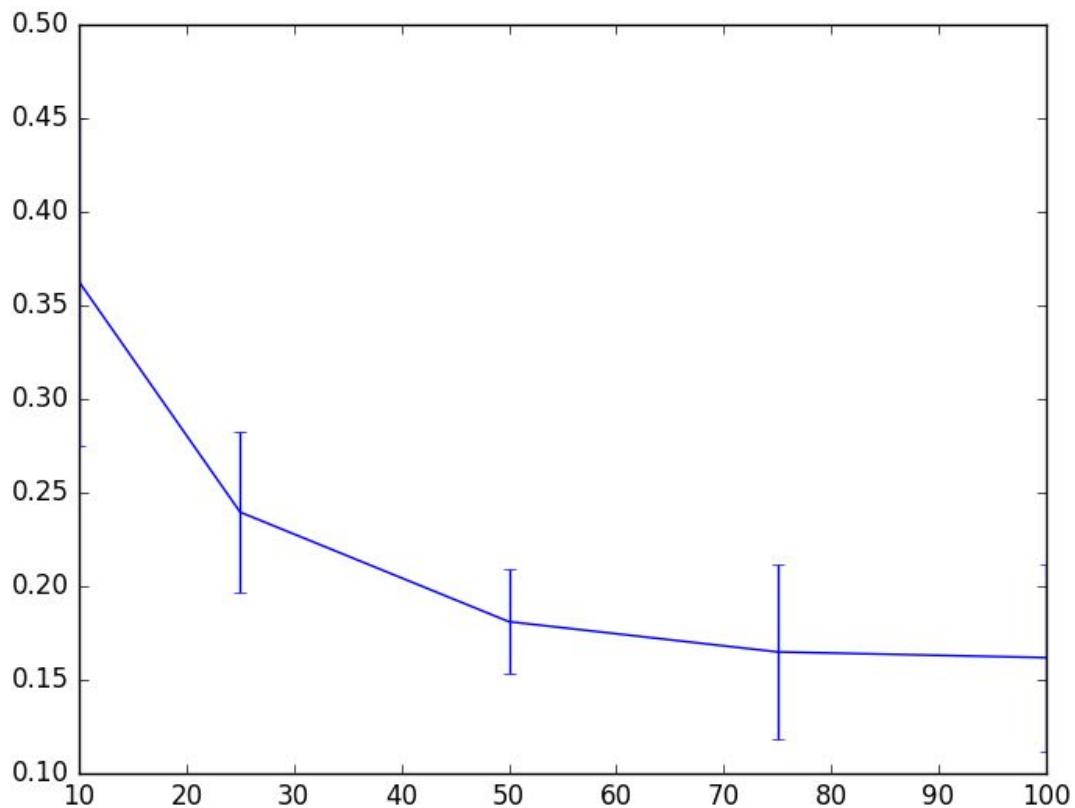
[ 0.25905292 0.1643454 0.1086351 0.09470752 0.11977716]

[ 0.15877437 0.09749304 0.08913649 0.07799443 0.07799443]]

Error standard deviation:

[ 0.05700174 0.03637098 0.01920996 0.01916953 0.03314646]

(ii)



Mean\_error\_rate:

[ 0.36267409 0.23955432 0.1810585 0.16490251 0.16183844]

Error\_matrix:(row index represents num\_split and column index is train\_percent)

```
[[ 0.4735376 0.26462396 0.19777159 0.17827298 0.16155989]
 [ 0.47632312 0.29805014 0.22284123 0.25069638 0.25626741]
 [ 0.32033426 0.22841226 0.18662953 0.17270195 0.14206128]
 [ 0.47075209 0.29805014 0.22005571 0.25069638 0.25626741]
 [ 0.40668524 0.2367688 0.13370474 0.11420613 0.09192201]
 [ 0.23398329 0.16155989 0.15041783 0.14206128 0.1448468 ]
 [ 0.23119777 0.16991643 0.15041783 0.13649025 0.14206128]
 [ 0.32590529 0.23955432 0.1810585 0.1448468 0.13927577]
 [ 0.3454039 0.25069638 0.18662953 0.13370474 0.1448468 ]
 [ 0.34261838 0.24791086 0.1810585 0.12534819 0.13927577]]
```

Error standard deviation:

[ 0.08747405 0.0432787 0.02779939 0.04669042 0.0500998 ]